

K-means clustering from scratch

Holly Capell

```
euc.dist <- function(x1, x2){
  sqrt(sum((x1 - x2)** 2))
}

kmeans <- function(nReps, myScatterInput, myClusterNum, maxIter){
  # Create empty list to be filled with sums of the final euclidean distances from the centroids
  euclidSums <- list()

  for (k in 1:nReps){
    # Create vector of cluster numbers
    clusterNums <- seq(1:myClusterNum)

    # Initialize cluster nums
    cluster_init <- sample(clusterNums, nrow(myScatterInput), replace = T)

    # Create cluster centroids by taking means of each group
    clusterCentroids <- bind_cols(myScatterInput, cluster = cluster_init) %>%
      group_by(cluster) %>%
      summarize_all(mean)

    cluster_new <- cluster_init

    count <- 0
    while (count < maxIter){
      cluster_old <- cluster_new

      # Create empty matrix to fill with distances
      dist_matrix <- matrix(0, nrow(myScatterInput), ncol=nrow((clusterCentroids)))

      # Turn cluster centroids into matrix
      cluster_matrix <- as.matrix(clusterCentroids)

      # Iterate over each cluster centroid to calculate euclidean distances with myScatterInput
      for (i in 1:nrow(clusterCentroids)){
        differences <- t(myScatterInput) - cluster_matrix[i,-1]

        eucDist <- t(sqrt(colSums(differences^2)))
        dist_matrix[,i] <- eucDist

      }

      # Find the minimum cluster distance for each observation in order to reassign clusters
      cluster_new <- (sapply(seq(nrow(dist_matrix)), function(i) {
        which.min(dist_matrix[i,])
      })))

      # Recalculate cluster centroids
      clusterCentroids <- bind_cols(myScatterInput, cluster = cluster_new) %>%
```

```

    group_by(cluster) %>%
    summarize_all(mean)

    # Break once cluster centroid assignments stop changing
    if (sum(cluster_new - cluster_old) == 0){
      break
    }
    count = count + 1
  }

  # Create data frame of final cluster assignments
  cluster_DF <- data.frame(cluster_new)
  colnames(cluster_DF) <- c("cluster")

  # Match up final cluster assignments to cluster means
  cluster_mat <- as.matrix(left_join(cluster_DF, clusterCentroids, by="cluster"))
  cluster_mat <- cluster_mat[,-1]

  # Compute the difference between observations and final centroid mean
  differences <- t(myScatterInput) - t(cluster_mat)

  # Compute euclidean distances and sum them
  eucDist <- t(sqrt(colSums(differences^2)))
  eucSum <- sum(eucDist)

  euclidSums[[k]] <- eucSum
}

# Find the minimum euclidean sum
minSum <- min(unlist(euclidSums))

if (ncol(myScatterInput) == 2){
  myScatterInput$cluster <- cluster_new
  colnames(myScatterInput) <- c("x1", "x2", "cluster")
  plot <- ggplot(myScatterInput, aes(x=x1, y=x2))+
    geom_point(aes(color=factor(cluster)))

  return(plot)
} else if (ncol(myScatterInput) == 3){
  myScatterInput$cluster <- cluster_new
  colnames(myScatterInput) <- c("x1", "x2", "x3", "cluster")
  colors <- c("red", "blue", "green", "orange", "yellow", "purple", "pink",
    "coral", "yellow", "navyblue", "turquoise", "olivedrab2")

  colors <- colors[as.numeric(myScatterInput$cluster)]
  plot <- scatterplot3d(myScatterInput[,1:3], color = colors)
  return(plot)
}
return (paste("Min Euclidean Sum: ",minSum))
}

# TEST DATA 1
set.seed(101)

```

```

myScatterInput <- data_frame(myCol_01 = runif(100000, -1, 1))
myClusterNum <- 2

kmeans(nReps = 5, myScatterInput = myScatterInput, myClusterNum = myClusterNum, maxIter = 10000)

## [1] "Min Euclidean Sum: 24862.2309460583"


microbenchmark(kmeans(1, myScatterInput, myClusterNum, 10000), times = 2)

## Unit: seconds
##              expr      min       lq     mean
## kmeans(1, myScatterInput, myClusterNum, 10000) 2.471587 2.471587 2.68203
##      median        uq      max neval
## 2.68203 2.892474 2.892474      2

# TEST DATA 1
set.seed(103)
myScatterInput <- data_frame(myCol_01 = runif(10000, -5, 20), myCol_02 = c(rnorm(3000, 20, 5), rnorm(5000, 40, 5), rnorm(2000, 0, 5)))
myClusterNum <- 3

kmeans(nReps = 5, myScatterInput = myScatterInput, myClusterNum = myClusterNum, maxIter = 10000)

```



```

# TEST DATA 2
set.seed(104)
myScatterInput <- data_frame(myCol_01 = c(rnorm(3000, 20, 20), rnorm(5000, -4, 2), rnorm(2000, 40, 2)),
myClusterNum <- 6

kmeans(nReps = 5, myScatterInput = myScatterInput, myClusterNum = myClusterNum, maxIter = 10000)

```

